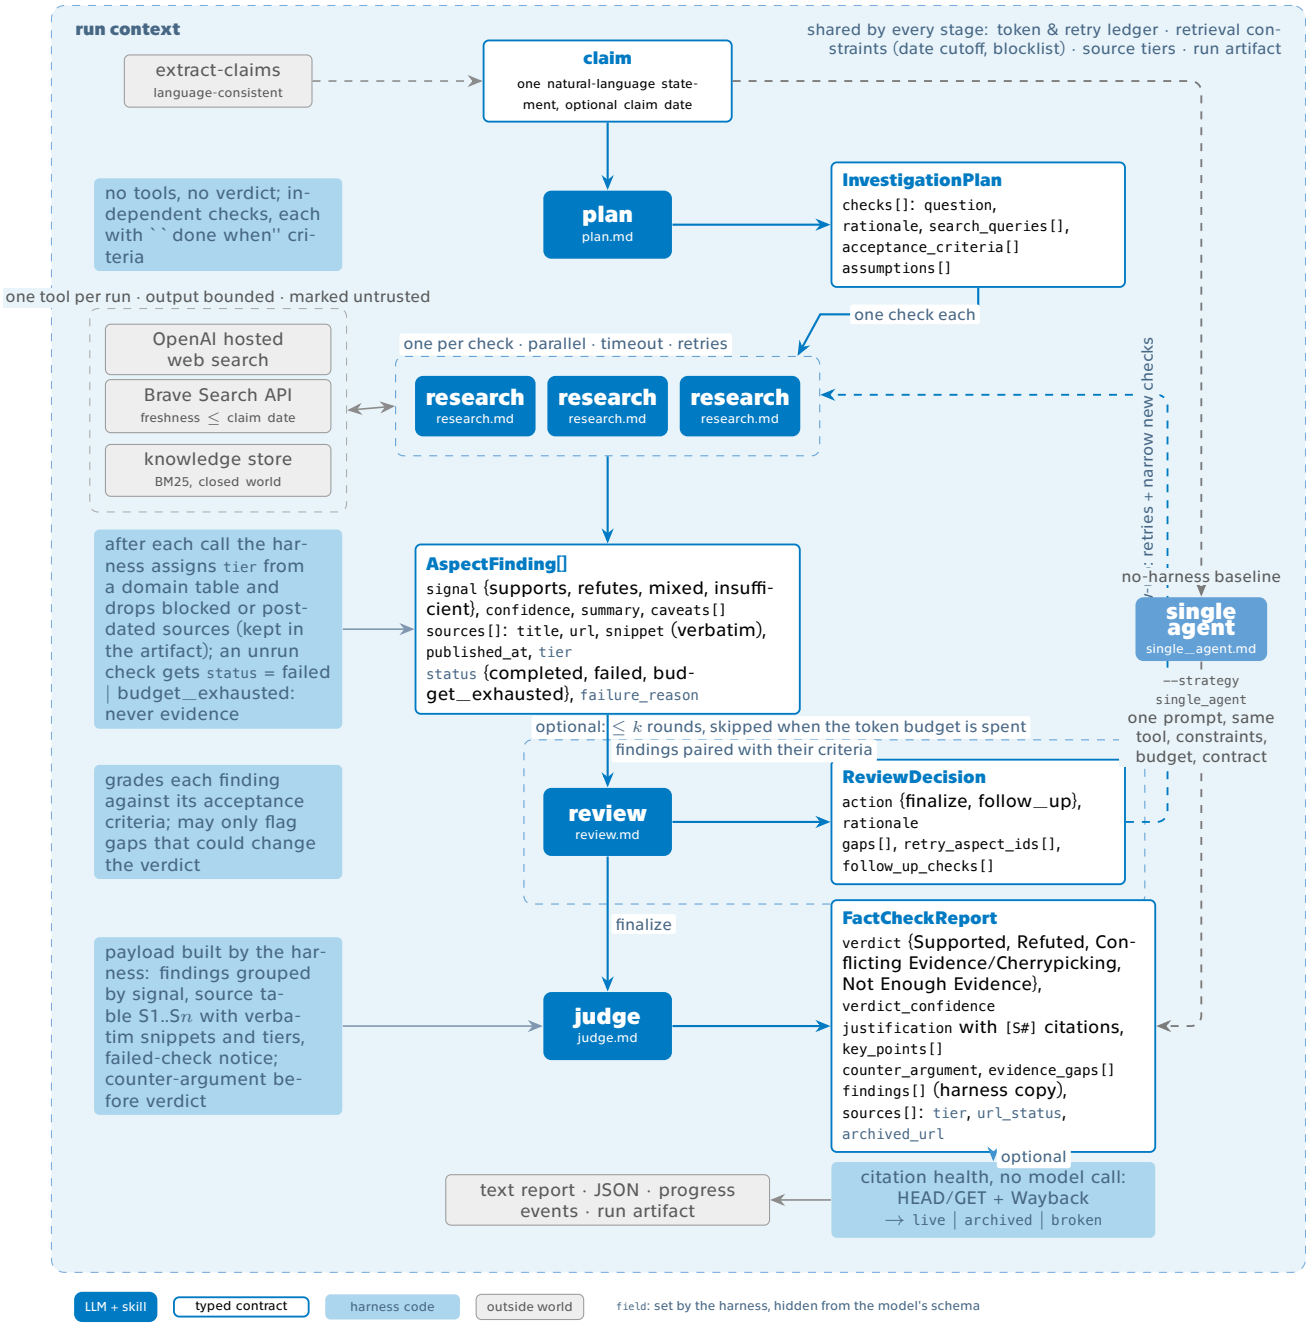
run context
shared by every stage: token & retry ledger · retrieval constraints (date cutoff, blocklist) · source tiers · run artifact
extract-claims
language-consistent
claim
one natural-language statement, optional claim date
plan
plan.md
no tools, no verdict; independent checks, each with ``done when'' criteria
InvestigationPlan
checks[]: question, rationale, search_queries[], acceptance_criteria[]
assumptions[]
one check each
one tool per run · output bounded · marked untrusted
OpenAI hosted web search
Brave Search API
freshness ≤ claim date
knowledge store
BM25, closed world
one per check · parallel · timeout · retries
research
research.md
research
research.md
research
research.md
after each call the harness assigns tier from a domain table and drops blocked or post-dated sources (kept in the artifact); an unrun check gets status = failed | budget_exhausted: never evidence
AspectFinding[]
signal {supports, refutes, mixed, insufficient}, confidence, summary, caveats[]
sources[]: title, url, snippet (verbatim), published_at, tier
status {completed, failed, budget_exhausted}, failure_reason
optional: ≤ k rounds, skipped when the token budget is spent
findings paired with their criteria
grades each finding against its acceptance criteria; may only flag gaps that could change the verdict
review
review.md
ReviewDecision
action {finalize, follow_up}, rationale
gaps[], retry_aspect_ids[], follow_up_checks[]
↩: retries + narrow new checks
finalize
payload built by the harness: findings grouped by signal, source table S1..Sn with verbatim snippets and tiers, failed-check notice; counter-argument before verdict
judge
judge.md
FactCheckReport
verdict {Supported, Refuted, Conflicting Evidence/Cherrypicking, Not Enough Evidence}, verdict_confidence
justification with [S#] citations, key_points[]
counter_argument, evidence_gaps[]
findings[] (harness copy), sources[]: tier, url_status, archived_url
optional
citation health, no model call: HEAD/GET + Wayback
→ live | archived | broken
text report · JSON · progress events · run artifact
no-harness baseline
single agent
single_agent.md
--strategy single_agent
one prompt, same tool, constraints, budget, contract
LLM + skill
typed contract
harness code
outside world
field: set by the harness, hidden from the model's schema